

WebP Container Specification

Introduction

WebP is an image format that uses either (i) the VP8 key frame encoding to compress image data in a lossy way, or (ii) the WebP lossless encoding (and possibly other encodings in the future). These encoding schemes should make it more efficient than currently used formats. It is optimized for fast image transfer over the network (e.g., for websites). The WebP format has feature parity (color profile, metadata, animation etc) with other formats as well. This document describes the structure of a WebP file.

The WebP container (i.e., RIFF container for WebP) allows feature support over and above the basic use case of WebP (i.e., a file containing a single image encoded as a VP8 key frame). The WebP container provides additional support for:

- **Lossless compression.** An image can be losslessly compressed, using the WebP Lossless Format.
- **Metadata.** An image may have metadata stored in EXIF or XMP formats.
- **Transparency.** An image may have transparency, i.e., an alpha channel.
- **Color Profile.** An image may have an embedded ICC profile as described by the [International Color Consortium](http://www.color.org/icc_specs2.xalter) (http://www.color.org/icc_specs2.xalter).
- **Animation.** An image may have multiple frames with pauses between them, making it an animation.

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC 2119](http://tools.ietf.org/html/rfc2119) (<http://tools.ietf.org/html/rfc2119>).

Bit numbering in chunk diagrams starts at 0 for the most significant bit ('MSB 0') as described in [RFC 1166](http://tools.ietf.org/html/rfc1166) (<http://tools.ietf.org/html/rfc1166>).

Naming

It is RECOMMENDED to use the following types when referring to the WebP container:

Container Format Name	WebP
Filename Extension	.webp
MIME-type	image/webp
Uniform Type Identifier	org.webmproject.webp

Terminology & Basics

A WebP file contains either a still image (i.e., an encoded matrix of pixels) or an animation (#animation). Optionally, it can also contain transparency information, color profile and metadata. In case we need to refer only to the matrix of pixels, we will call it the *canvas* of the image.

Below are additional terms used throughout this document:

Reader/Writer

Code that reads WebP files is referred to as a *reader*, while code that writes them is referred to as a *writer*.

uint16

A 16-bit, little-endian, unsigned integer.

uint24

A 24-bit, little-endian, unsigned integer.

uint32

A 32-bit, little-endian, unsigned integer.

FourCC

A *FourCC* (four-character code) is a *uint32* created by concatenating four ASCII characters in little-endian order.

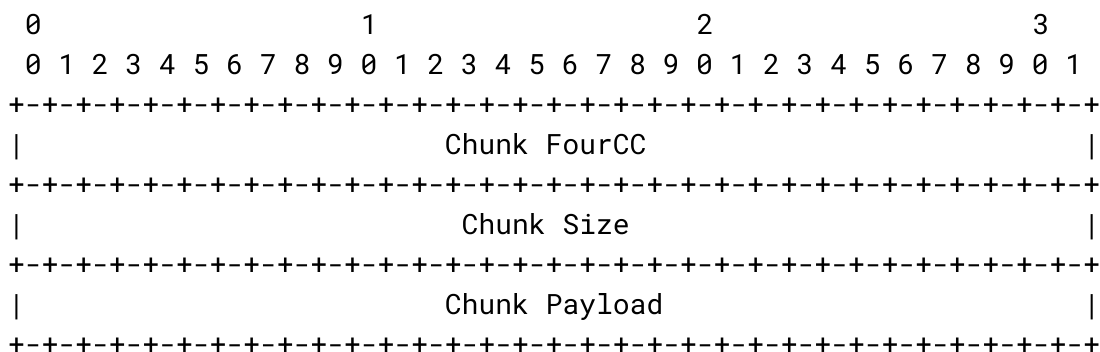
1-based

An unsigned integer field storing values offset by -1. e.g., Such a field would store value 25 as 24.

RIFF File Format

The WebP file format is based on the RIFF (resource interchange file format) document format.

The basic element of a RIFF file is a *chunk*. It consists of:



Chunk FourCC: 32 bits

ASCII four-character code used for chunk identification.

Chunk Size: 32 bits (*uint32*)

The size of the chunk not including this field, the chunk identifier or padding.

Chunk Payload: *Chunk Size* bytes

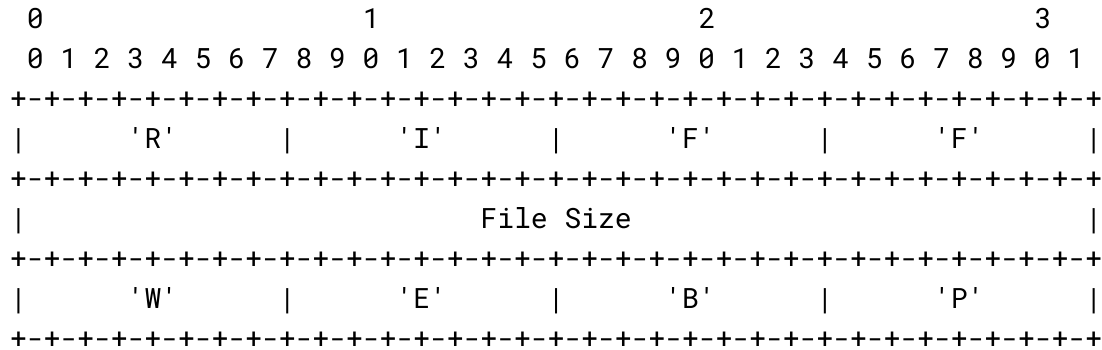
The data payload. If *Chunk Size* is odd, a single padding byte -- that SHOULD be 0 -- is added.

ChunkHeader('ABCD')

This is used to describe the *FourCC* and *Chunk Size* header of individual chunks, where 'ABCD' is the FourCC for the chunk. This element's size is 8 bytes.

Note: RIFF has a convention that all-uppercase chunk FourCCs are standard chunks that apply to any RIFF file format, while FourCCs specific to a file format are all lowercase. WebP does not follow this convention.

WebP File Header



'RIFF': 32 bits

The ASCII characters 'R' 'I' 'F' 'F'.

File Size: 32 bits (*uint32*)

The size of the file in bytes starting at offset 8. The maximum value of this field is 2^{32} minus 10 bytes and thus the size of the whole file is at most 4GiB minus 2 bytes.

'WEBP': 32 bits

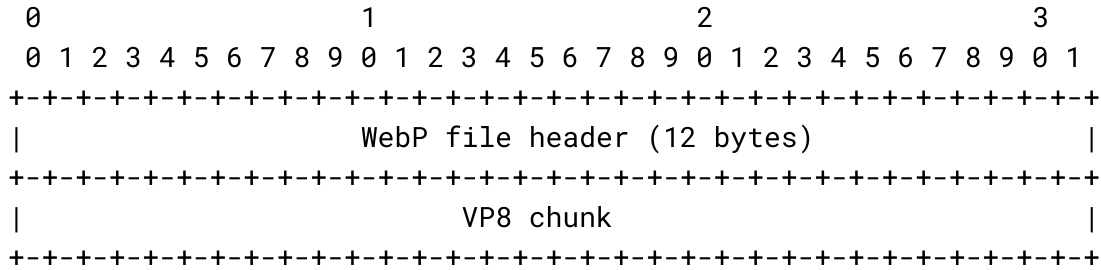
The ASCII characters 'W' 'E' 'B' 'P'.

A WebP file MUST begin with a RIFF header with the FourCC 'WEBP'. The file size in the header is the total size of the chunks that follow plus 4 bytes for the 'WEBP' FourCC. The file SHOULD NOT contain anything after it. As the size of any chunk is even, the size given by the RIFF header is also even. The contents of individual chunks will be described in the following sections.

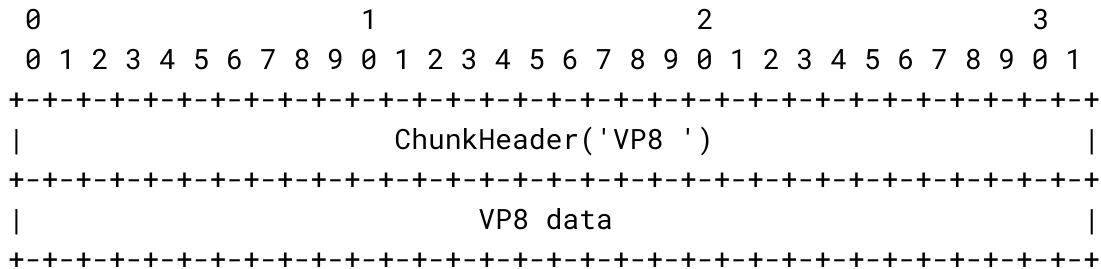
Simple File Format (Lossy)

This layout SHOULD be used if the image requires *lossy* encoding and does not require transparency or other advanced features provided by the extended format. Files with this layout are smaller and supported by older software.

Simple WebP (lossy) file format:



VP8 chunk:



VP8 data: *Chunk Size* bytes

VP8 bitstream data.

The VP8 bitstream format specification can be found at [VP8 Data Format and Decoding Guide](http://tools.ietf.org/html/rfc6386) (<http://tools.ietf.org/html/rfc6386>). Note that the VP8 frame header contains the VP8 frame width and height. That is assumed to be the width and height of the canvas.

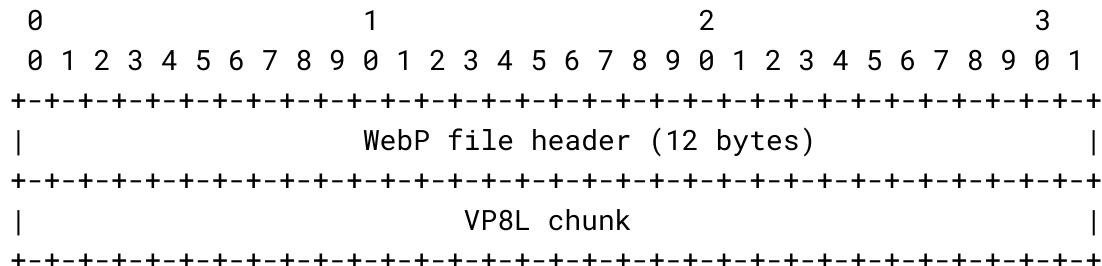
The VP8 specification describes how to decode the image into Y'CbCr format. To convert to RGB, Rec. 601 SHOULD be used.

Simple File Format (Lossless)

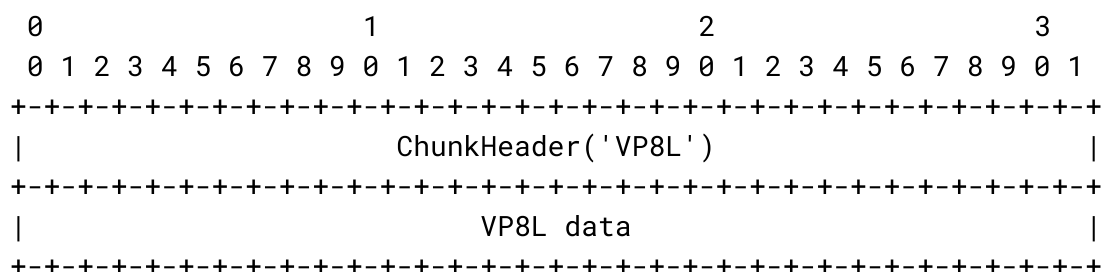
Note: Older readers may not support files using the lossless format.

This layout SHOULD be used if the image requires *lossless* encoding (with an optional transparency channel) and does not require advanced features provided by the extended format.

Simple WebP (lossless) file format:



VP8L chunk:



VP8L data: *Chunk Size* bytes

VP8L bitstream data.

The current specification of the VP8L bitstream can be found at [WebP Lossless Bitstream Format](https://developers.google.com/speed/webp/docs/webp_lossless_bitstream_specification) (https://developers.google.com/speed/webp/docs/webp_lossless_bitstream_specification). Note that the VP8L header contains the VP8L image width and height. That is assumed to be the width and height of the canvas.

Extended File Format

Note: Older readers may not support files using the extended format.

An extended format file consists of:

- A 'VP8X' chunk with information about features used in the file.
- An optional 'ICCP' chunk with color profile.
- An optional 'ANIM' chunk with animation control data.
- Image data.
- An optional 'EXIF' chunk with EXIF metadata.
- An optional 'XMP' chunk with XMP metadata.

For a *still image*, the *image data* consists of a single frame, which is made up of:

- An optional alpha subchunk (#alpha).
- A bitstream subchunk (#bitstream_vp8vp8l).

For an *animated image*, the *image data* consists of multiple frames. More details about frames can be found in the Animation (#animation) section.

All chunks SHOULD be placed in the same order as listed above. If a chunk appears in the wrong place, the file is invalid, but readers MAY parse the file, ignoring the chunks that come too late.

Rationale: Setting the order of chunks should allow quicker file parsing. For example, if an 'ALPH' chunk does not appear in its required position, a decoder can choose to stop searching for it. The rule of ignoring late chunks should make programs that need to do a full search give the same results as the ones stopping early.

Extended WebP file header:

```

0                               1                               2                               3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               WebP file header (12 bytes)                               |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               ChunkHeader('VP8X')                               |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|Rsv|I|L|E|X|A|R|                               Reserved                               |

```

```

+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|           Canvas Width Minus One           |           ...
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
... Canvas Height Minus One   |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

Reserved (Rsv): 2 bits

SHOULD be 0.

ICC profile (I): 1 bit

Set if the file contains an ICC profile.

Alpha (L): 1 bit

Set if any of the frames of the image contain transparency information ("alpha").

EXIF metadata (E): 1 bit

Set if the file contains EXIF metadata.

XMP metadata (X): 1 bit

Set if the file contains XMP metadata.

Animation (A): 1 bit

Set if this is an animated image. Data in 'ANIM' and 'ANMF' chunks should be used to control the animation.

Reserved (R): 1 bit

SHOULD be 0.

Reserved: 24 bits

SHOULD be 0.

Canvas Width Minus One: 24 bits

1-based width of the canvas in pixels. The actual canvas width is '1 + Canvas Width Minus One'

Canvas Height Minus One: 24 bits

1-based height of the canvas in pixels. The actual canvas height is '1 + Canvas Height Minus One'

The product of *Canvas Width* and *Canvas Height* MUST be at most $2^{32} - 1$.

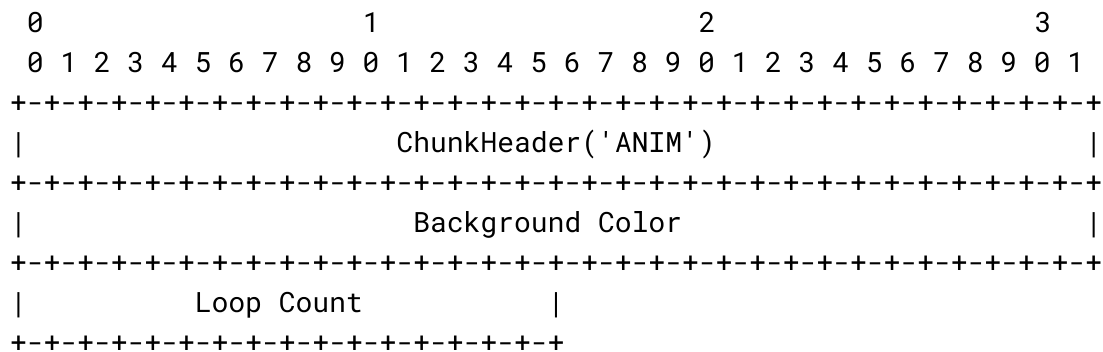
Future specifications MAY add more fields.

Animation

An animation is controlled by ANIM and ANMF chunks.

ANIM Chunk:

For an animated image, this chunk contains the *global parameters* of the animation.



Background Color: 32 bits (*uint32*)

The default background color of the canvas in [Blue, Green, Red, Alpha] byte order. This color MAY be used to fill the unused space on the canvas around the frames, as well as the transparent pixels of the first frame. Background color is also used when disposal method is 1.

Note:

- Background color MAY contain a transparency value (alpha), even if the *Alpha* flag in VP8X chunk (*#extended_header*) is unset.
- Viewer applications SHOULD treat the background color value as a hint, and are not required to use it.

- The canvas is cleared at the start of each loop. The background color MAY be used to achieve this.

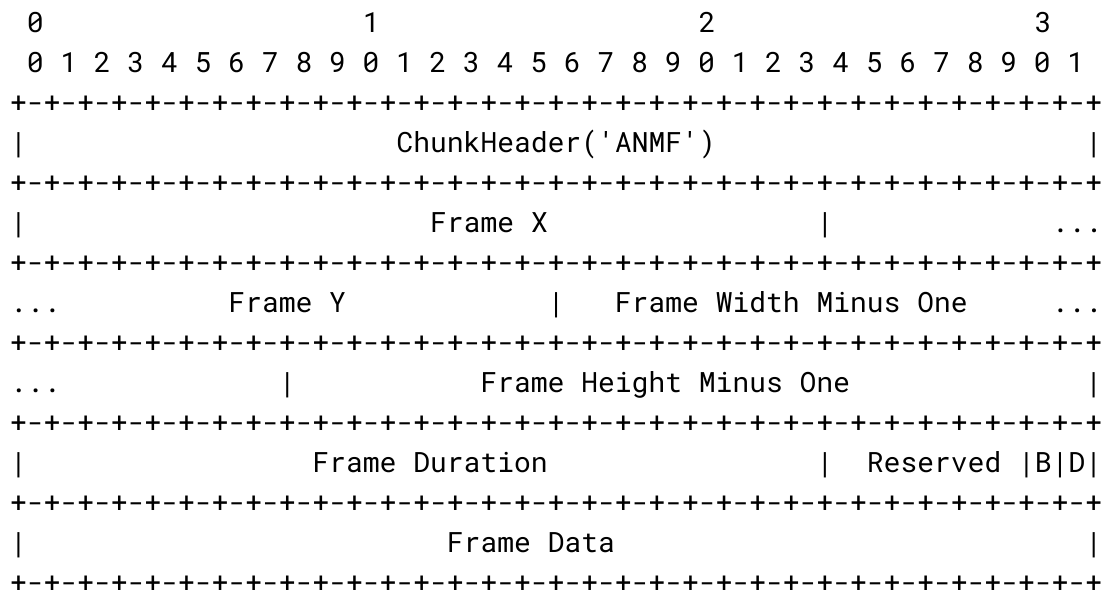
Loop Count: 16 bits (*uint16*)

The number of times to loop the animation. 0 means infinitely.

This chunk MUST appear if the *Animation* flag in the VP8X chunk is set. If the *Animation* flag is not set and this chunk is present, it SHOULD be ignored.

ANMF chunk:

For animated images, this chunk contains information about a *single* frame. If the *Animation flag* is not set, then this chunk SHOULD NOT be present.



Frame X: 24 bits (*uint24*)

The X coordinate of the upper left corner of the frame is $\text{Frame X} * 2$

Frame Y: 24 bits (*uint24*)

The Y coordinate of the upper left corner of the frame is `Frame Y * 2`

Frame Width Minus One: 24 bits (*uint24*)

The 1-based width of the frame. The frame width is 1 + Frame Width Minus One

Frame Height Minus One: 24 bits (*uint24*)

The *1-based* height of the frame. The frame height is $1 + \text{Frame Height Minus One}$

Frame Duration: 24 bits (*uint24*)

The time to wait before displaying the next frame, in 1 millisecond units. Note the interpretation of frame duration of 0 (and often ≤ 10) is implementation defined. Many tools and browsers assign a minimum duration similar to GIF.

Reserved: 6 bits

SHOULD be 0.

Blending method (B): 1 bit

Indicates how transparent pixels of *the current frame* are to be blended with corresponding pixels of the previous canvas:

- 0: Use alpha blending. After disposing of the previous frame, render the current frame on the canvas using alpha-blending (see below). If the current frame does not have an alpha channel, assume alpha value of 255, effectively replacing the rectangle.
- 1: Do not blend. After disposing of the previous frame, render the current frame on the canvas by overwriting the rectangle covered by the current frame.

Disposal method (D): 1 bit

Indicates how *the current frame* is to be treated after it has been displayed (before rendering the next frame) on the canvas:

- 0: Do not dispose. Leave the canvas as is.
- 1: Dispose to background color. Fill the *rectangle* on the canvas covered by the *current frame* with background color specified in the ANIM chunk (`#anim_chunk`).

Notes:

- The frame disposal only applies to the *frame rectangle*, that is, the rectangle defined by *Frame X*, *Frame Y*, *frame width* and *frame height*. It may or may not cover the whole canvas.

- **Alpha-blending:** Given that each of the R, G, B and A channels is 8-bit, and the RGB channels are *not premultiplied* by alpha, the formula for blending 'dst' onto 'src' is:

```
blend.A = src.A + dst.A * (1 - src.A / 255)
if blend.A = 0 then
    blend.RGB = 0
else
    blend.RGB = (src.RGB * src.A +
                 dst.RGB * dst.A * (1 - src.A / 255)) / blend.A
```

- Alpha-blending SHOULD be done in linear color space, by taking into account the color profile (#color_profile) of the image. If the color profile is not present, sRGB is to be assumed. (Note that sRGB also needs to be linearized due to a gamma of ~ 2.2).

Frame Data: *Chunk Size* - 16 bytes

- A bitstream subchunk (#bitstream_vp8vp8l) for the frame.

Note: The 'ANMF' payload, *Frame Data* above, consists of individual *padded* chunks as described by the RIFF file format (#riff_file_format).

Alpha

```

0                                     1                                     2                                     3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+-+-+-+-+-+
|                                     ChunkHeader('ALPH')                                     |
+-+-+-+-+-+-+
|Rsv| P | F | C | Alpha Bitstream... |
+-+-+-+-+-+-+

```

Reserved (Rsv): 2 bits

SHOULD be 0.

Pre-processing (P): 2 bits

These INFORMATIVE bits are used to signal the pre-processing that has been performed during compression. The decoder can use this information to e.g. dither the values or smooth the gradients prior to display.

- 0: no pre-processing
- 1: level reduction

Filtering method (F): 2 bits

The filtering method used:

- 0: None.
- 1: Horizontal filter.
- 2: Vertical filter.
- 3: Gradient filter.

For each pixel, filtering is performed using the following calculations. Assume the alpha values surrounding the current X position are labeled as:

```

C | B |
---+---+
A | X |

```

We seek to compute the alpha value at position X. First, a prediction is made depending on the filtering method:

- Method 0: predictor = 0
- Method 1: predictor = A
- Method 2: predictor = B
- Method 3: predictor = clip(A + B - C)

where clip(v) is equal to:

- 0 if $v < 0$
- 255 if $v > 255$

- v otherwise

The final value is derived by adding the decompressed value X to the predictor and using modulo-256 arithmetic to wrap the $[256-511]$ range into the $[0-255]$ one:

$\alpha = (\text{predictor} + X) \% 256$

There are special cases for left-most and top-most pixel positions:

- Top-left value at location $(0,0)$ uses 0 as predictor value. Otherwise,
- For horizontal or gradient filtering methods, the left-most pixels at location $(0, y)$ are predicted using the location $(0, y-1)$ just above.
- For vertical or gradient filtering methods, the top-most pixels at location $(x, 0)$ are predicted using the location $(x-1, 0)$ on the left.

Decoders are not required to use this information in any specified way.

Compression method (C): 2 bits

The compression method used:

- 0: No compression.
- 1: Compressed using the WebP lossless format.

Alpha bitstream: *Chunk Size* - 1 bytes

Encoded alpha bitstream.

This optional chunk contains encoded alpha data for this frame. A frame containing a 'VP8L' chunk SHOULD NOT contain this chunk.

Rationale: The transparency information is already part of the 'VP8L' chunk.

The alpha channel data is stored as uncompressed raw data (when compression method is '0') or compressed using the lossless format (when the compression method is '1').

- Raw data: consists of a byte sequence of length $\text{width} * \text{height}$, containing all the 8-bit transparency values in scan order.
- Lossless format compression: the byte sequence is a compressed image-stream (as described in the [WebP Lossless Bitstream Format](#))

(https://developers.google.com/speed/webp/docs/webp_lossless_bitstream_specification)) of implicit dimension width x height. That is, this image-stream does NOT contain any headers describing the image dimension.

Rationale: the dimension is already known from other sources, so storing it again would be redundant and error-prone.

Once the image-stream is decoded into ARGB color values, following the process described in the lossless format specification, the transparency information must be extracted from the *green* channel of the ARGB quadruplet.

Rationale: the green channel is allowed extra transformation steps in the specification – unlike the other channels – that can improve compression.

Bitstream (VP8/VP8L)

This chunk contains compressed bitstream data for a single frame.

A bitstream chunk may be either (i) a VP8 chunk, using "VP8 " (note the significant fourth-character space) as its tag or (ii) a VP8L chunk, using "VP8L" as its tag.

The formats of VP8 and VP8L chunks are as described in sections [Simple File Format \(Lossy\)](#). (#simple_file_format_lossy) and [Simple File Format \(Lossless\)](#). (#simple_file_format_lossless) respectively.

Color profile

[illegible]

Color Profile: *Chunk Size* bytes

ICC profile.

This chunk **MUST** appear before the image data.

There **SHOULD** be at most one such chunk. If there are more such chunks, readers **MAY** ignore all except the first one. See the [ICC Specification](http://www.color.org/icc_specs2.xalter) (http://www.color.org/icc_specs2.xalter) for details.

If this chunk is not present, sRGB **SHOULD** be assumed.

Metadata

Metadata can be stored in 'EXIF' or 'XMP' chunks.

There **SHOULD** be at most one chunk of each type ('EXIF' and 'XMP'). If there are more such chunks, readers **MAY** ignore all except the first one. Also, a file may possibly contain both 'EXIF' and 'XMP' chunks.

The chunks are defined as follows:

EXIF chunk:

```

0                               1                               2                               3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               ChunkHeader( ' EXIF ' )                               |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               EXIF Metadata                                         |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

EXIF Metadata: *Chunk Size* bytes

image metadata in EXIF format.

XMP chunk:

```

0                               1                               2                               3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|                               ChunkHeader( ' XMP ' )                               |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```


| XMP Metadata |

XMP Metadata: *Chunk Size* bytes

image metadata in XMP format.

Additional guidance about handling metadata can be found in the Metadata Working Group's Guidelines for Handling Metadata (http://www.metadataworkinggroup.org/pdf/mwq_guidance.pdf).

Displaying an *animated image* canvas MUST be equivalent to the following pseudocode:

```

assert VP8X.flags.hasAnimation
canvas ← new image of size VP8X.canvasWidth x VP8X.canvasHeight with
    background color ANIM.background_color.
loop_count ← ANIM.loopCount
dispose_method ← ANIM.disposeMethod
if loop_count == 0:
    loop_count = ∞
frame_params ← nil
assert next chunk in image_data is ANMF
for loop = 0..loop_count - 1
    clear canvas to ANIM.background_color or application defined color
    until eof or non-ANMF chunk
        frame_params.frameX = Frame X
        frame_params.frameY = Frame Y
        frame_params.frameWidth = Frame Width Minus One + 1
        frame_params.frameHeight = Frame Height Minus One + 1
        frame_params.frameDuration = Frame Duration
        frame_right = frame_params.frameX + frame_params.frameWidth
        frame_bottom = frame_params.frameY + frame_params.frameHeight
        assert VP8X.canvasWidth ≥ frame_right
        assert VP8X.canvasHeight ≥ frame_bottom
        for subchunk in 'Frame Data':
            if subchunk.tag == "ALPH":
                assert alpha subchunks not found in 'Frame Data' earlier
                frame_params.alpha = alpha_data
            else if subchunk.tag == "VP8 " OR subchunk.tag == "VP8L":
                assert bitstream subchunks not found in 'Frame Data' earlier
                frame_params.bitstream = bitstream_data
    render frame with frame_params.alpha and frame_params.bitstream on
    canvas with top-left corner at (frame_params.frameX,
    frame_params.frameY), using dispose method dispose_method.

```

```
canvas contains the decoded image.  
Show the contents of the canvas for  
    frame_params.frameDuration * 1ms.
```

Example File Layouts

A lossy encoded image with alpha may look as follows:

```
RIFF/WEBP  
+- VP8X (descriptions of features used)  
+- ALPH (alpha bitstream)  
+- VP8 (bitstream)
```

A losslessly encoded image may look as follows:

```
RIFF/WEBP  
+- VP8X (descriptions of features used)  
+- VP8L (lossless bitstream)
```

A lossless image with ICC profile and XMP metadata may look as follows:

```
RIFF/WEBP  
+- VP8X (descriptions of features used)  
+- ICCP (color profile)  
+- VP8L (lossless bitstream)  
+- XMP (metadata)
```

An animated image with EXIF metadata may look as follows:

```
RIFF/WEBP  
+- VP8X (descriptions of features used)  
+- ANIM (global animation parameters)  
+- ANMF (frame1 parameters + data)  
+- ANMF (frame2 parameters + data)
```

- + - ANMF (frame3 parameters + data)
- + - ANMF (frame4 parameters + data)
- + - EXIF (metadata)

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2021-04-14 UTC.